

Stephanie Patricia Anshell

sanshell@ucsd.edu | (925) 348-7120 | [linkedin.com/in/stephaniepatriciaanshell](https://www.linkedin.com/in/stephaniepatriciaanshell) | github.com/stephaniepatriciaans | stephaniepatriciaans.github.io/portfolio

AI/ML engineer and incoming Data Science graduate (UCSD, 2026) building with the Claude API, GenAI prompt engineering, and agentic LLM workflows across 3 internships spanning real estate, supply chain, and urban analytics. Proven ability to ship predictive models and automation pipelines fast, translate model outputs into stakeholder decisions, and structure ambiguous business problems into measurable analytical questions. Proven ability to ship real tools fast, translate model outputs into stakeholder decisions, and structure ambiguous business problems into measurable analytical questions. Proficient in Python, SQL, JavaScript, scikit-learn, TensorFlow, and Tableau. Fluent in English and Indonesian.

EDUCATION

UNIVERSITY OF CALIFORNIA SAN DIEGO

Bachelor of Science in Data Science and Minor in Finance | Cumulative GPA: 3.6

La Jolla, California, USA

Expected 2026

UNIVERSITY OF CALIFORNIA BERKELEY EXTENSION

Certificate Data Analytics, Data Analytics Bootcamp

Berkeley, California, USA

November 2023

SKILLS

AI / LLM: GenAI, Claude API, Prompt Engineering, Agentic Pipelines, RAG, Tool Use / Function Calling, Multi-step LLM Workflows

Languages: Python, SQL, JavaScript, MATLAB, C++

ML / Data: scikit-learn, TensorFlow, pandas, NumPy, SciPy, GeoPandas, PySAL, Matplotlib, Seaborn, PyTorch, Spark

Methods: Hypothesis-Driven Problem Solving, ML, Deep Learning, NLP, Predictive Modeling, Time-Series Forecasting, A/B Testing, Feature Engineering, ETL, Data Visualization, Business Strategy, Business Transformation

Tools: Tableau, Streamlit, Jupyter, Git, Linux, ROS, AWS, Microsoft 365 (Excel, SharePoint, Teams)

Automation: API Integration (REST), Data Pipeline Automation, Workflow Automation, Cross-Functional Stakeholder Collaboration

TECHNICAL EXPERIENCE

UC SAN DIEGO TECHNICAL TEAMS (Triton AI Member & Autonomous Racing Team)

ROS, Linux, Python, Robotics, Autonomous Systems, Simulation | Autonomous Systems / Robotics Contributor

La Jolla, California

January 2026 - Present

- Advance applied AI, ML, robotics, and autonomous systems initiatives through Triton AI and the UCSD Autonomous Racing Team, contributing to technical discussions, workshops & project development across AI/ML, LLM workflows, ROS & simulation environments.
- Supported a 2nd-place finish at a Purdue University competition by documenting 3+ technical subsystems of a ROS/Linux autonomous vehicle project, improving reproducible onboarding workflows for new contributors.

EXPERIENCE

IDX Exchange (Real Estate Data Science Platform)

Data Scientist Intern

Remote, USA

September 2025 - December 2025

- Engineered a client-facing ML pipeline (XGBoost, Random Forest, Linear Regression) integrated via Cotality's Trestle API, achieving $R^2=0.89$ and ~6% MDAPE across 5+ home-price segments, then translated results into executive decision briefs that shaped pricing strategy.
- Cut analysis turnaround time 40% by designing a standardized, repeatable workflow across 5+ experiments (EDA → feature engineering → tuning → evaluation), enabling faster, more reliable client-ready outputs.

GAF (Authorized distributor of GAF Roofing materials in Indonesia, focused on supply chain operations)

Data Scientist Intern

Jakarta, Indonesia

January 2024 - September 2024

- Structured an ambiguous "reduce stock-outs" business question into a measurable forecasting problem; built ARIMA and Decision Tree models on 9 months of procurement data, lifting forecast accuracy 15% and cutting stock-outs 20% for purchasing stakeholders.
- Saved 3 hrs/week and eliminated manual errors in KPI reporting by building Tableau dashboards and automated data pipelines tracking 8+ metrics across procurement and inventory.

STAL Corporation (A printing supplier for Indonesian hospitals and clinics)

Accounting & Data Analyst Intern

Jakarta, Indonesia

Jun 2023- Aug 2023 - Hybrid/Remote

- Drove a 15% cut in procurement costs by identifying hidden over-ordering patterns through time-series analysis, informing corrective purchasing policy across finance and operations.
- Recovered 4 hrs/week in manual reconciliation time by designing automated expense-tracking and budget-planning workflows, improving financial reporting precision across monthly cycles.

PROJECTS

Loan Charge Offs | Python, scikit-learn, TensorFlow | [GitHub : github.com/stephaniepatriciaans/Project-Loan-Charge-Offs](https://github.com/stephaniepatriciaans/Project-Loan-Charge-Offs)

- Built an ETL pipeline on FRED data (390 rows, 1991–present) and compared 4 ML methods; Random Forest hit 93.91% accuracy, with threshold analysis for business precision-recall tradeoffs.
- Constructed an ETL pipeline converting monthly/quarterly FRED data into a unified modeling dataset with automated quality checks, enabling reproducible credit analysis across economic cycles.

Urban Transit & Crime | GeoPandas, PySAL | [GitHub : github.com/stephaniepatriciaans/Urban-Transit-and-Crime](https://github.com/stephaniepatriciaans/Urban-Transit-and-Crime)

- Led a team of 5 to integrate 1M+ records across 4 public datasets (CTA, MTA, NYPD, CPD), directing the build of spatial regression models (GeoPandas, PySAL, spreg) that detected statistically significant transit-crime clusters across Chicago and NYC.

TrueFork | Python, scikit-learn, Pandas, JavaScript | [GitHub : github.com/stephaniepatriciaans/TrueFork](https://github.com/stephaniepatriciaans/TrueFork) | [Live Site :](https://stephaniepatriciaans.github.io/TrueFork/)

- Built an NLP pipeline on 80K+ Food.com recipe-interaction records to classify review sentiment into pros/cons, reducing simulated user decision time ~50%; shipped as a live public site, demonstrating end-to-end AI tool deployment.
- Validated feature significance at $p<0.05$ via A/B testing and statistical modeling; documented full pipeline as a reproducible GitHub workflow with interactive data visualizations.

Adult Obesity Prevalence | Python, scikit-learn, pandas | github.com/stephaniepatriciaans/adult-obesity-prevalence

- Independently built and owned the full regression pipeline (OLS, Ridge, Lasso) on CDC PLACES data across 2,299 counties and 7 health/social-need predictors, achieving $R^2=0.628$ via 5-fold CV, as 1 of 3 team members, coordinating with teammates on data sourcing and final report/poster deliverables.